

1.2.1 : EDA: Advanced Feature Extraction.

```
1 import warnings
2 warnings.filterwarnings("ignore")
3 import numpy as np
4 import pandas as pd
5 import seaborn as sns
6 import matplotlib.pyplot as plt
7 from subprocess import check_output
8 %matplotlib inline
9 import plotly.offline as py
10 py.init_notebook_mode(connected=True)
11 import plotly.graph_objs as go
12 import plotly.tools as tls
13 import os
14 import gc
15
16 import re
17 from nltk.corpus import stopwords
18 import distance
19 from nltk.stem import PorterStemmer
20 from bs4 import BeautifulSoup
21 import re
22 from nltk.corpus import stopwords
23 # This package is used for finding longest common subsequence between two strings
24 # you can write your own dp code for this
25 import distance
26 from nltk.stem import PorterStemmer
27 from bs4 import BeautifulSoup
28 from fuzzywuzzy import fuzz
29 from sklearn.manifold import TSNE
30 # Import the Required Lib packages for WORD-Cloud generation
31 # https://stackoverflow.com/questions/45625434/how-to-install-wordcloud-in-python3-6
32 from wordcloud import WordCloud, STOPWORDS
33 from os import path
34 from PIL import Image
```

```

1 #https://stackoverflow.com/questions/12468179/unicodedecodeerror-utf8-codec-cant-decode-byte-0x9c
2 if os.path.isfile('df_fe_without_preprocessing_train.csv'):
3     df = pd.read_csv("df_fe_without_preprocessing_train.csv",encoding='latin-1')
4     df = df.fillna('')
5     df.head()
6 else:
7     print("get df_fe_without_preprocessing_train.csv from drive or run the previous notebook")

```

```
1 df.head(2)
```

	id	qid1	qid2	question1	question2	is_duplicate	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n
0	0	1	2	What is the step by step guide to invest in sh...	What is the step by step guide to invest in sh...	0	1	1	66	57	14	12
1	1	3	4	What is the story of Kohinoor (Koh-i-Noor) Dia...	What would happen if the Indian government sto...	0	4	1	51	88	8	13

3.4 Preprocessing of Text

- Preprocessing:
 - Removing html tags
 - Removing Punctuations
 - Performing stemming
 - Removing Stopwords
 - Expanding contractions etc.

```
1  # To get the results in 4 decemal points
2  SAFE_DIV = 0.0001
3
4  STOP_WORDS = stopwords.words("english")
5
6
7  def preprocess(x):
8      x = str(x).lower()
9      x = x.replace(",000,000", "m").replace(",000", "k").replace("'", "").replace('"', '')\
10         .replace("won't", "will not").replace("cannot", "can not").replace("can't", "
11         .replace("n't", " not").replace("what's", "what is").replace("it's", "it is")
12         .replace("'ve", " have").replace("i'm", "i am").replace("'re", " are")\
13         .replace("he's", "he is").replace("she's", "she is").replace("'s", " own")\
14         .replace("%", " percent ").replace("₹", " rupee ").replace("$", " dollar ")\
15         .replace("€", " euro ").replace("'ll", " will")
16      x = re.sub(r"([0-9]+)000000", r"\1m", x)
17      x = re.sub(r"([0-9]+)000", r"\1k", x)
18
19
20      porter = PorterStemmer()
21      pattern = re.compile('\W')
22
23      if type(x) == type(''):
24          x = re.sub(pattern, ' ', x)
25
26
27      if type(x) == type(''):
28          x = porter.stem(x)
29          example1 = BeautifulSoup(x)
30          x = example1.get_text()
31
32
33      return x
34
```

- Function to Compute and get the features : With 2 parameters of Question 1 and Question 2

3.5 Advanced Feature Extraction (NLP and Fuzzy Features)

Definition:

- **Token**: You get a token by splitting sentence a space
- **Stop_Word** : stop words as per NLTK.
- **Word** : A token that is not a stop_word

Features:

- **cwc_min** : Ratio of common_word_count to min length of word count of Q1 and Q2
$$\text{cwc_min} = \text{common_word_count} / (\min(\text{len}(\text{q1_words}), \text{len}(\text{q2_words})))$$
- **cwc_max** : Ratio of common_word_count to max length of word count of Q1 and Q2
$$\text{cwc_max} = \text{common_word_count} / (\max(\text{len}(\text{q1_words}), \text{len}(\text{q2_words})))$$
- **csc_min** : Ratio of common_stop_count to min length of stop count of Q1 and Q2
$$\text{csc_min} = \text{common_stop_count} / (\min(\text{len}(\text{q1_stops}), \text{len}(\text{q2_stops})))$$
- **csc_max** : Ratio of common_stop_count to max length of stop count of Q1 and Q2
$$\text{csc_max} = \text{common_stop_count} / (\max(\text{len}(\text{q1_stops}), \text{len}(\text{q2_stops})))$$
- **ctc_min** : Ratio of common_token_count to min length of token count of Q1 and Q2
$$\text{ctc_min} = \text{common_token_count} / (\min(\text{len}(\text{q1_tokens}), \text{len}(\text{q2_tokens})))$$
- **ctc_max** : Ratio of common_token_count to max length of token count of Q1 and Q2
$$\text{ctc_max} = \text{common_token_count} / (\max(\text{len}(\text{q1_tokens}), \text{len}(\text{q2_tokens})))$$
- **last_word_eq** : Check if First word of both questions is equal or not
$$\text{last_word_eq} = \text{int}(\text{q1_tokens}[-1] == \text{q2_tokens}[-1])$$
- **first_word_eq** : Check if First word of both questions is equal or not
$$\text{first_word_eq} = \text{int}(\text{q1_tokens}[0] == \text{q2_tokens}[0])$$

- **abs_len_diff** : Abs. length difference
 $\text{abs_len_diff} = \text{abs}(\text{len}(q1_tokens) - \text{len}(q2_tokens))$
- **mean_len** : Average Token Length of both Questions
 $\text{mean_len} = (\text{len}(q1_tokens) + \text{len}(q2_tokens))/2$
- **fuzz_ratio** : <https://github.com/seatgeek/fuzzywuzzy#usage> (<https://github.com/seatgeek/fuzzywuzzy#usage>)
<http://chairnerd.seatgeek.com/fuzzywuzzy-fuzzy-string-matching-in-python/>
(<http://chairnerd.seatgeek.com/fuzzywuzzy-fuzzy-string-matching-in-python/>)
- **fuzz_partial_ratio** : <https://github.com/seatgeek/fuzzywuzzy#usage>
(<https://github.com/seatgeek/fuzzywuzzy#usage>) <http://chairnerd.seatgeek.com/fuzzywuzzy-fuzzy-string-matching-in-python/>
(<http://chairnerd.seatgeek.com/fuzzywuzzy-fuzzy-string-matching-in-python/>)
- **token_sort_ratio** : <https://github.com/seatgeek/fuzzywuzzy#usage>
(<https://github.com/seatgeek/fuzzywuzzy#usage>) <http://chairnerd.seatgeek.com/fuzzywuzzy-fuzzy-string-matching-in-python/>
(<http://chairnerd.seatgeek.com/fuzzywuzzy-fuzzy-string-matching-in-python/>)
- **token_set_ratio** : <https://github.com/seatgeek/fuzzywuzzy#usage>
(<https://github.com/seatgeek/fuzzywuzzy#usage>) <http://chairnerd.seatgeek.com/fuzzywuzzy-fuzzy-string-matching-in-python/>
(<http://chairnerd.seatgeek.com/fuzzywuzzy-fuzzy-string-matching-in-python/>)
- **longest_substr_ratio** : Ratio of length longest common substring to min length of token count of Q1 and Q2
 $\text{longest_substr_ratio} = \text{len}(\text{longest common substring}) / (\min(\text{len}(q1_tokens), \text{len}(q2_tokens)))$

```
1 def get_token_features(q1, q2):
2     token_features = [0.0]*10
3
4     # Converting the Sentence into Tokens:
5     q1_tokens = q1.split()
6     q2_tokens = q2.split()
7
8     if len(q1_tokens) == 0 or len(q2_tokens) == 0:
9         return token_features
10
11     # Get the non-stopwords in Questions
12     q1_words = set([word for word in q1_tokens if word not in STOP_WORDS])
13     q2_words = set([word for word in q2_tokens if word not in STOP_WORDS])
14
15     #Get the stopwords in Questions
16     q1_stops = set([word for word in q1_tokens if word in STOP_WORDS])
17     q2_stops = set([word for word in q2_tokens if word in STOP_WORDS])
18
19     # Get the common non-stopwords from Question pair
20     common_word_count = len(q1_words.intersection(q2_words))
21
22     # Get the common stopwords from Question pair
23     common_stop_count = len(q1_stops.intersection(q2_stops))
24
25     # Get the common Tokens from Question pair
26     common_token_count = len(set(q1_tokens).intersection(set(q2_tokens)))
27
28     token_features[0] = common_word_count / (min(len(q1_words), len(q2_words)) + SAFE_DIV)
29     token_features[1] = common_word_count / (max(len(q1_words), len(q2_words)) + SAFE_DIV)
30     token_features[2] = common_stop_count / (min(len(q1_stops), len(q2_stops)) + SAFE_DIV)
31     token_features[3] = common_stop_count / (max(len(q1_stops), len(q2_stops)) + SAFE_DIV)
32     token_features[4] = common_token_count / (min(len(q1_tokens), len(q2_tokens)) + SAFE_DIV)
33     token_features[5] = common_token_count / (max(len(q1_tokens), len(q2_tokens)) + SAFE_DIV)
34
```

```
35     # Last word of both question is same or not
36     token_features[6] = int(q1_tokens[-1] == q2_tokens[-1])
37
38     # First word of both question is same or not
39     token_features[7] = int(q1_tokens[0] == q2_tokens[0])
40
41     token_features[8] = abs(len(q1_tokens) - len(q2_tokens))
42
43     #Average Token Length of both Questions
44     token_features[9] = (len(q1_tokens) + len(q2_tokens))/2
45     return token_features
46
47     # get the Longest Common sub string
48
49     def get_longest_substr_ratio(a, b):
50         strs = list(distance.lcs substrings(a, b))
51         if len(strs) == 0:
52             return 0
53         else:
54             return len(strs[0]) / (min(len(a), len(b)) + 1)
55
56     def extract_features(df):
57         # preprocessing each question
58         df["question1"] = df["question1"].fillna("").apply(preprocess)
59         df["question2"] = df["question2"].fillna("").apply(preprocess)
60
61         print("token features...")
62
63         # Merging Features with dataset
64
65         token_features = df.apply(lambda x: get_token_features(x["question1"], x["question2"]), axis=1)
66
67         df["cwc_min"] = list(map(lambda x: x[0], token_features))
68         df["cwc_max"] = list(map(lambda x: x[1], token_features))
69         df["csc_min"] = list(map(lambda x: x[2], token_features))
```



```
70 df["csc_max"] = list(map(lambda x: x[3], token_features))
71 df["ctc_min"] = list(map(lambda x: x[4], token_features))
72 df["ctc_max"] = list(map(lambda x: x[5], token_features))
73 df["last_word_eq"] = list(map(lambda x: x[6], token_features))
74 df["first_word_eq"] = list(map(lambda x: x[7], token_features))
75 df["abs_len_diff"] = list(map(lambda x: x[8], token_features))
76 df["mean_len"] = list(map(lambda x: x[9], token_features))
77
78 #Computing Fuzzy Features and Merging with Dataset
79
80 # do read this blog: http://chairnerd.seatgeek.com/fuzzywuzzy-fuzzy-string-matching-in-python/
81 # https://stackoverflow.com/questions/31806695/when-to-use-which-fuzz-function-to-compare-2-strings
82 # https://github.com/seatgeek/fuzzywuzzy
83 print("fuzzy features..")
84
85 df["token_set_ratio"] = df.apply(lambda x: fuzz.token_set_ratio(x["question1"], x["question2"]), axis=1)
86 # The token sort approach involves tokenizing the string in question, sorting the tokens alphabetically
87 # then joining them back into a string We then compare the transformed strings with a simple ratio()
88 df["token_sort_ratio"] = df.apply(lambda x: fuzz.token_sort_ratio(x["question1"], x["question2"]), axis=1)
89 df["fuzz_ratio"] = df.apply(lambda x: fuzz.QRatio(x["question1"], x["question2"]), axis=1)
90 df["fuzz_partial_ratio"] = df.apply(lambda x: fuzz.partial_ratio(x["question1"], x["question2"]), axis=1)
91 df["longest_substr_ratio"] = df.apply(lambda x: get_longest_substr_ratio(x["question1"], x["question2"]), axis=1)
92 return df
```

```

1 if os.path.isfile('nlp_features_train.csv'):
2     df = pd.read_csv("nlp_features_train.csv",encoding='latin-1')
3     df.fillna('')
4 else:
5     print("Extracting features for train:")
6     df = pd.read_csv("train.csv")
7     df = extract_features(df)
8     df.to_csv("nlp_features_train.csv", index=False)
9 df.head(2)

```

	id	qid1	qid2	question1	question2	is_duplicate	cwc_min	cwc_max	csc_min	csc_max	...	ctc_max
0	0	1	2	what is the step by step guide to invest in sh...	what is the step by step guide to invest in sh...	0	0.999980	0.833319	0.999983	0.999983	...	0.785709
1	1	3	4	what is the story of kohinoor koh i noor dia...	what would happen if the indian government sto...	0	0.799984	0.399996	0.749981	0.599988	...	0.466664

2 rows x 21 columns

3.5.1 Analysis of extracted features

3.5.1.1 Plotting Word clouds

- Creating Word Cloud of Duplicates and Non-Duplicates Question pairs
- We can observe the most frequent occurring words

```
1 df_duplicate = df[df['is_duplicate'] == 1]
2 dfp_nonduplicate = df[df['is_duplicate'] == 0]
3
4 # Converting 2d array of q1 and q2 and flatten the array: Like {{1,2},{3,4}} to {1,2,3,4}
5 p = np.dstack([df_duplicate["question1"], df_duplicate["question2"]]).flatten()
6 n = np.dstack([dfp_nonduplicate["question1"], dfp_nonduplicate["question2"]]).flatten()
7
8 print ("Number of data points in class 1 (duplicate pairs) :",len(p))
9 print ("Number of data points in class 0 (non duplicate pairs) :",len(n))
10
11 #Saving the np array into a text file
12 np.savetxt('train_p.txt', p, delimiter=' ', fmt='%s')
13 np.savetxt('train_n.txt', n, delimiter=' ', fmt='%s')
```

Number of data points in class 1 (duplicate pairs) : 298526


Number of data points in class 0 (non duplicate pairs) : 510054

```
1 # reading the text files and removing the Stop Words:
2 d = path.dirname('.')
3
4 textp_w = open(path.join(d, 'train_p.txt')).read()
5 textn_w = open(path.join(d, 'train_n.txt')).read()
6 stopwords = set(STOPWORDS)
7 stopwords.add("said")
8 stopwords.add("br")
9 stopwords.add(" ")
10 stopwords.remove("not")
11
12 stopwords.remove("no")
13 #stopwords.remove("good")
14 #stopwords.remove("Love")
15 stopwords.remove("like")
16 #stopwords.remove("best")
17 #stopwords.remove("!")
18 print ("Total number of words in duplicate pair questions :",len(textp_w))
19 print ("Total number of words in non duplicate pair questions :",len(textn_w))
```


Total number of words in duplicate pair questions : 16109886
Total number of words in non duplicate pair questions : 33193130

__ Word Clouds generated from duplicate pair question's text __

Word Cloud for Duplicate Question pairs

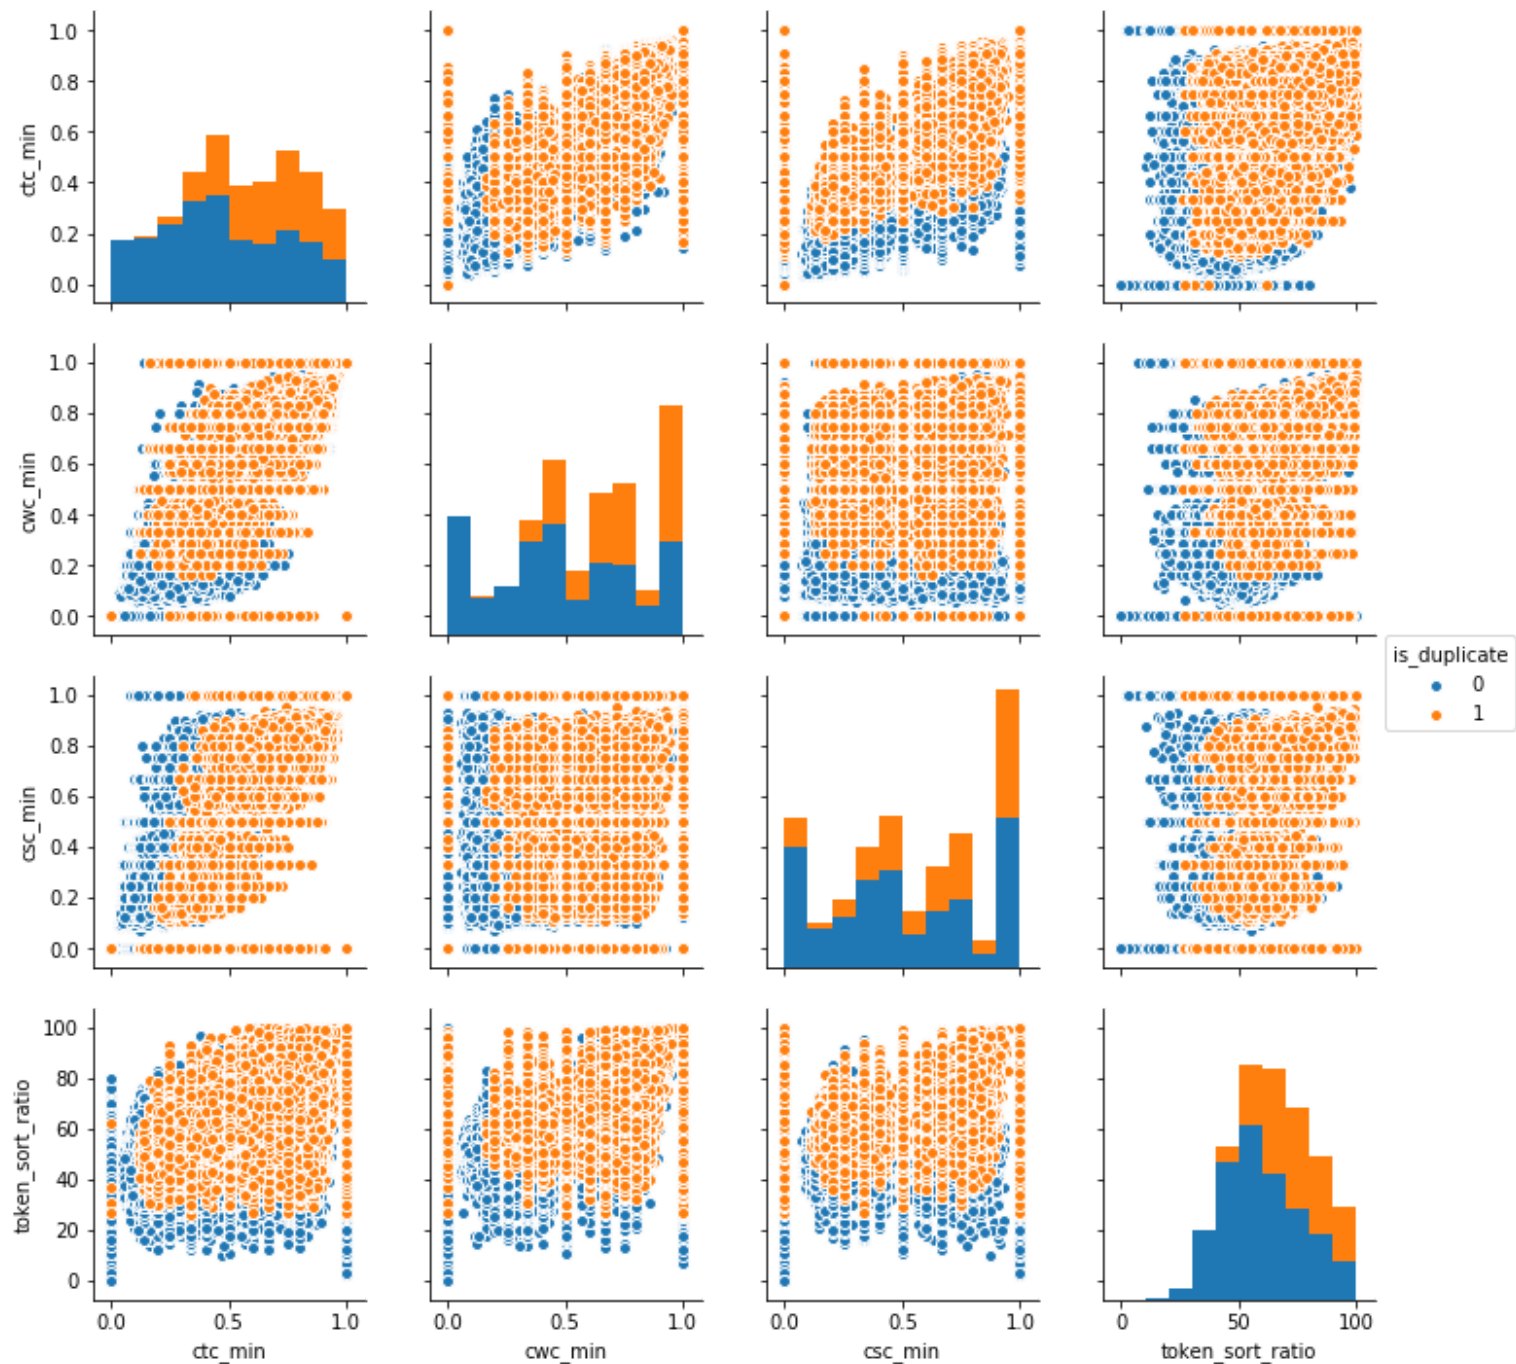


Word Cloud for non-Duplicate Question pairs:



http://localhost:8888/notebooks/Documents/Applied%20AI%20assignments/Quora%20Question%20pair%20similarity%20problem/2.Quora_Preprocessing.ipynb


```
1 n = df.shape[0]
2 sns.pairplot(df[['ctc_min', 'cwc_min', 'csc_min', 'token_sort_ratio', 'is_duplicate']][0:n], hue='is_du
3 plt.show()
```



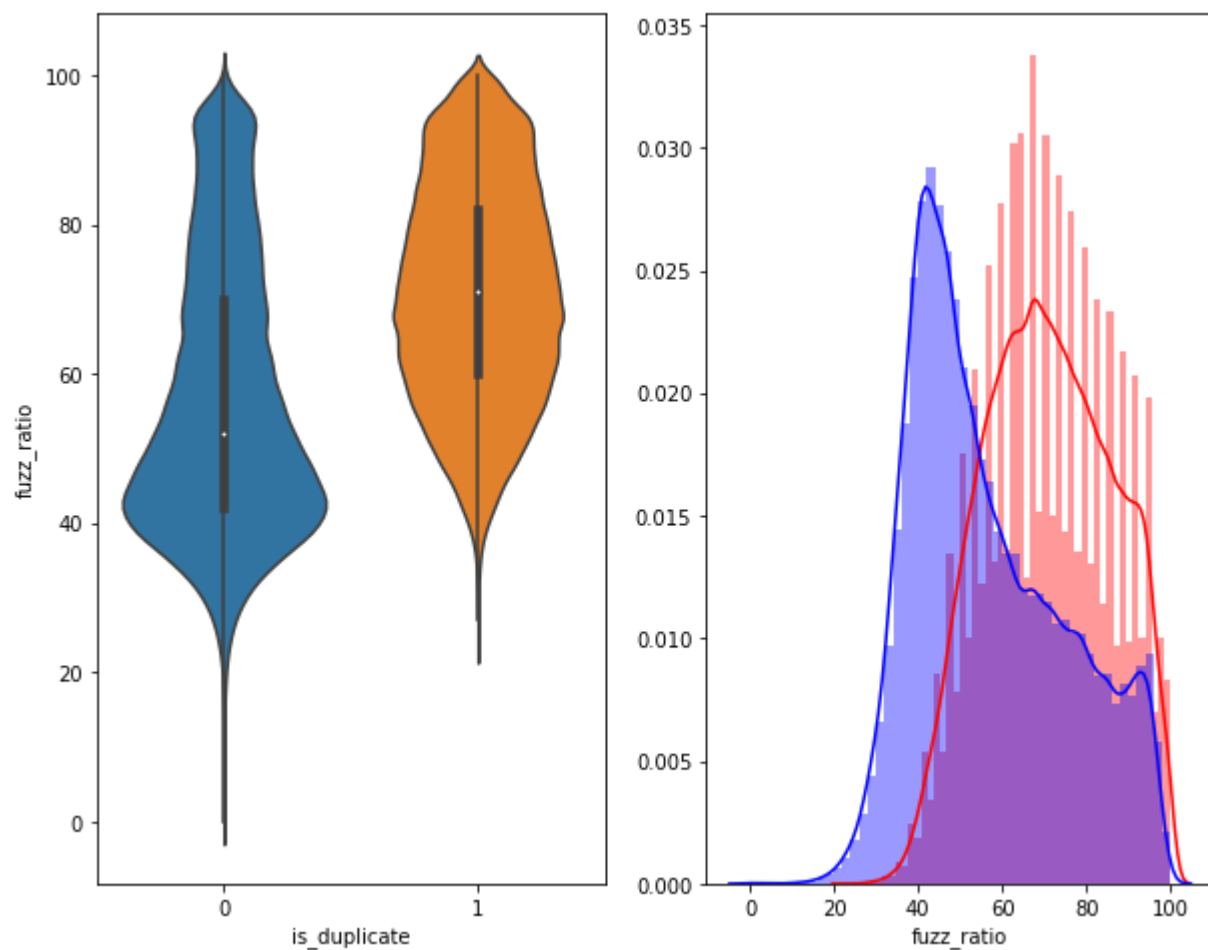

```

1 # Distribution of the token_sort_ratio
2 plt.figure(figsize=(10, 8))
3
4 plt.subplot(1,2,1)
5 sns.violinplot(x = 'is_duplicate', y = 'token_sort_ratio', data = df[0:] , )
6
7 plt.subplot(1,2,2)
8 sns.distplot(df[df['is_duplicate'] == 1.0]['token_sort_ratio'][0:] , label = "1", color = 'red')
9 sns.distplot(df[df['is_duplicate'] == 0.0]['token_sort_ratio'][0:] , label = "0" , color = 'blue' )
10 plt.show()

```



```
1 plt.figure(figsize=(10, 8))
2
3 plt.subplot(1,2,1)
4 sns.violinplot(x = 'is_duplicate', y = 'fuzz_ratio', data = df[0:] , )
5
6 plt.subplot(1,2,2)
7 sns.distplot(df[df['is_duplicate'] == 1.0]['fuzz_ratio'][0:] , label = "1", color = 'red')
8 sns.distplot(df[df['is_duplicate'] == 0.0]['fuzz_ratio'][0:] , label = "0" , color = 'blue' )
9 plt.show()
```



3.5.2 Visualization


```
1 # Using TSNE for Dimentionality reduction for 15 Features(Generated after cleaning the data) to 3 dimen
2
3 from sklearn.preprocessing import MinMaxScaler
4
5 dfp_subsampled = df[0:5000]
6 X = MinMaxScaler().fit_transform(dfp_subsampled[['cwc_min', 'cwc_max', 'csc_min', 'csc_max' , 'ctc_min']
7 y = dfp_subsampled['is_duplicate'].values
```

```
1 tsne2d = TSNE(  
2     n_components=2,  
3     init='random', # pca  
4     random_state=101,  
5     method='barnes_hut',  
6     n_iter=1000,  
7     verbose=2,  
8     angle=0.5  
9 ).fit_transform(X)
```

```
[t-SNE] Computing 91 nearest neighbors...  
[t-SNE] Indexed 5000 samples in 0.011s...  
[t-SNE] Computed neighbors for 5000 samples in 0.912s...  
[t-SNE] Computed conditional probabilities for sample 1000 / 5000  
[t-SNE] Computed conditional probabilities for sample 2000 / 5000  
[t-SNE] Computed conditional probabilities for sample 3000 / 5000  
[t-SNE] Computed conditional probabilities for sample 4000 / 5000  
[t-SNE] Computed conditional probabilities for sample 5000 / 5000  
[t-SNE] Mean sigma: 0.116557  
[t-SNE] Computed conditional probabilities in 0.433s  
[t-SNE] Iteration 50: error = 80.9244080, gradient norm = 0.0428133 (50 iterations in 13.099s)  
[t-SNE] Iteration 100: error = 70.3858795, gradient norm = 0.0100968 (50 iterations in 9.067s)  
[t-SNE] Iteration 150: error = 68.6138382, gradient norm = 0.0058392 (50 iterations in 9.602s)  
[t-SNE] Iteration 200: error = 67.7700119, gradient norm = 0.0036596 (50 iterations in 9.121s)  
[t-SNE] Iteration 250: error = 67.2725067, gradient norm = 0.0034962 (50 iterations in 11.305s)  
[t-SNE] KL divergence after 250 iterations with early exaggeration: 67.272507  
[t-SNE] Iteration 300: error = 1.7737305, gradient norm = 0.0011918 (50 iterations in 8.289s)  
[t-SNE] Iteration 350: error = 1.3720417, gradient norm = 0.0004822 (50 iterations in 10.526s)  
[t-SNE] Iteration 400: error = 1.2039998, gradient norm = 0.0002768 (50 iterations in 9.600s)  
[t-SNE] Iteration 450: error = 1.1133438, gradient norm = 0.0001881 (50 iterations in 11.827s)  
[t-SNE] Iteration 500: error = 1.0579143, gradient norm = 0.0001434 (50 iterations in 8.941s)  
[t-SNE] Iteration 550: error = 1.0221983, gradient norm = 0.0001164 (50 iterations in 11.092s)  
[t-SNE] Iteration 600: error = 0.9987167, gradient norm = 0.0001039 (50 iterations in 11.467s)  
[t-SNE] Iteration 650: error = 0.9831534, gradient norm = 0.0000938 (50 iterations in 11.799s)  
[t-SNE] Iteration 700: error = 0.9722011, gradient norm = 0.0000858 (50 iterations in 12.028s)  
[t-SNE] Iteration 750: error = 0.9643636, gradient norm = 0.0000799 (50 iterations in 12.120s)  
[t-SNE] Iteration 800: error = 0.9584482, gradient norm = 0.0000785 (50 iterations in 11.867s)  
[t-SNE] Iteration 850: error = 0.9538348, gradient norm = 0.0000739 (50 iterations in 11.461s)  
[t-SNE] Iteration 900: error = 0.9496906, gradient norm = 0.0000712 (50 iterations in 11.023s)  
[t-SNE] Iteration 950: error = 0.9463405, gradient norm = 0.0000673 (50 iterations in 11.755s)
```

```
[t-SNE] Iteration 1000: error = 0.9432716, gradient norm = 0.0000662 (50 iterations in 11.493s)  
[t-SNE] Error after 1000 iterations: 0.943272
```

```
1 df = pd.DataFrame({'x':tsne2d[:,0], 'y':tsne2d[:,1] , 'label':y})
2
3 # draw the plot in appropriate place in the grid
4 sns.lmplot(data=df, x='x', y='y', hue='label', fit_reg=False, size=8,palette="Set1",markers=['s','o'])
5 plt.title("perplexity : {} and max_iter : {}".format(30, 1000))
6 plt.show()
```



```
1 from sklearn.manifold import TSNE
2 tsne3d = TSNE(
3     n_components=3,
4     init='random', # pca
5     random_state=101,
6     method='barnes_hut',
7     n_iter=1000,
8     verbose=2,
9     angle=0.5
10 ).fit_transform(X)
```

```
[t-SNE] Computing 91 nearest neighbors...
[t-SNE] Indexed 5000 samples in 0.010s...
[t-SNE] Computed neighbors for 5000 samples in 0.935s...
[t-SNE] Computed conditional probabilities for sample 1000 / 5000
[t-SNE] Computed conditional probabilities for sample 2000 / 5000
[t-SNE] Computed conditional probabilities for sample 3000 / 5000
[t-SNE] Computed conditional probabilities for sample 4000 / 5000
[t-SNE] Computed conditional probabilities for sample 5000 / 5000
[t-SNE] Mean sigma: 0.116557
[t-SNE] Computed conditional probabilities in 0.363s
[t-SNE] Iteration 50: error = 77.7944183, gradient norm = 0.1014017 (50 iterations in 34.931s)
[t-SNE] Iteration 100: error = 69.2682266, gradient norm = 0.0248657 (50 iterations in 15.147s)
[t-SNE] Iteration 150: error = 67.7877655, gradient norm = 0.0150941 (50 iterations in 13.761s)
[t-SNE] Iteration 200: error = 67.1991119, gradient norm = 0.0126559 (50 iterations in 13.425s)
[t-SNE] Iteration 250: error = 66.8560715, gradient norm = 0.0074975 (50 iterations in 12.904s)
[t-SNE] KL divergence after 250 iterations with early exaggeration: 66.856071
[t-SNE] Iteration 300: error = 1.2356015, gradient norm = 0.0007033 (50 iterations in 13.302s)
[t-SNE] Iteration 350: error = 0.9948602, gradient norm = 0.0001997 (50 iterations in 18.898s)
[t-SNE] Iteration 400: error = 0.9168936, gradient norm = 0.0001430 (50 iterations in 13.397s)
[t-SNE] Iteration 450: error = 0.8863022, gradient norm = 0.0000975 (50 iterations in 16.379s)
[t-SNE] Iteration 500: error = 0.8681002, gradient norm = 0.0000854 (50 iterations in 17.791s)
[t-SNE] Iteration 550: error = 0.8564141, gradient norm = 0.0000694 (50 iterations in 17.060s)
[t-SNE] Iteration 600: error = 0.8470711, gradient norm = 0.0000640 (50 iterations in 15.454s)
[t-SNE] Iteration 650: error = 0.8389117, gradient norm = 0.0000561 (50 iterations in 17.562s)
[t-SNE] Iteration 700: error = 0.8325295, gradient norm = 0.0000529 (50 iterations in 13.443s)
[t-SNE] Iteration 750: error = 0.8268463, gradient norm = 0.0000528 (50 iterations in 17.981s)
[t-SNE] Iteration 800: error = 0.8219477, gradient norm = 0.0000477 (50 iterations in 17.448s)
[t-SNE] Iteration 850: error = 0.8180174, gradient norm = 0.0000490 (50 iterations in 18.376s)
[t-SNE] Iteration 900: error = 0.8150476, gradient norm = 0.0000456 (50 iterations in 17.778s)
```



```
[t-SNE] Iteration 950: error = 0.8122067, gradient norm = 0.0000472 (50 iterations in 16.983s)
[t-SNE] Iteration 1000: error = 0.8095787, gradient norm = 0.0000489 (50 iterations in 18.581s)
[t-SNE] Error after 1000 iterations: 0.809579
```

```
1 trace1 = go.Scatter3d(  
2     x=tsne3d[:,0],  
3     y=tsne3d[:,1],  
4     z=tsne3d[:,2],  
5     mode='markers',  
6     marker=dict(  
7         sizemode='diameter',  
8         color = y,  
9         colorscale = 'Portland',  
10        colorbar = dict(title = 'duplicate'),  
11        line=dict(color='rgb(255, 255, 255)'),  
12        opacity=0.75  
13    )  
14 )  
15  
16 data=[trace1]  
17 layout=dict(height=800, width=800, title='3d embedding with engineered features')  
18 fig=dict(data=data, layout=layout)  
19 py.iplot(fig, filename='3DBubble')
```

